



Single-Source Shortest paths in large-scale dynamic networks

By:

Arrij Fawwad 22i-0755

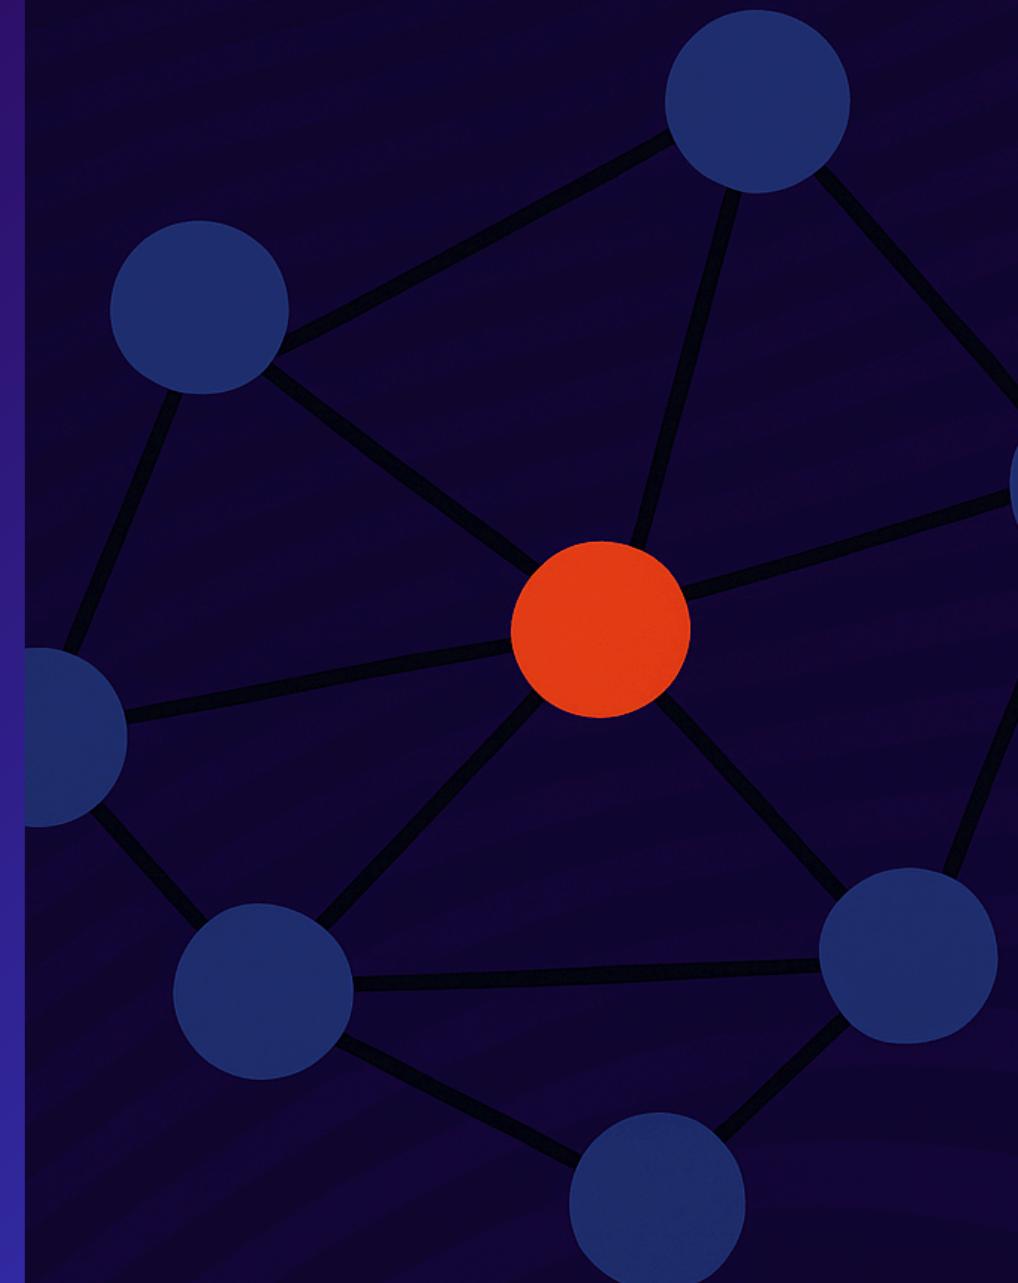
Hamna Arshad 22i-1098

Maryam Masood 22i-1169

INTRODUCTION

- SSSP: Shortest paths from one source to all nodes
- Real-world graphs change over time
- Recomputing SSSP is slow and inefficient
- Existing methods often platform-specific
- Proposed: fast, parallel SSSP update framework

KEY CONTRIBUTION



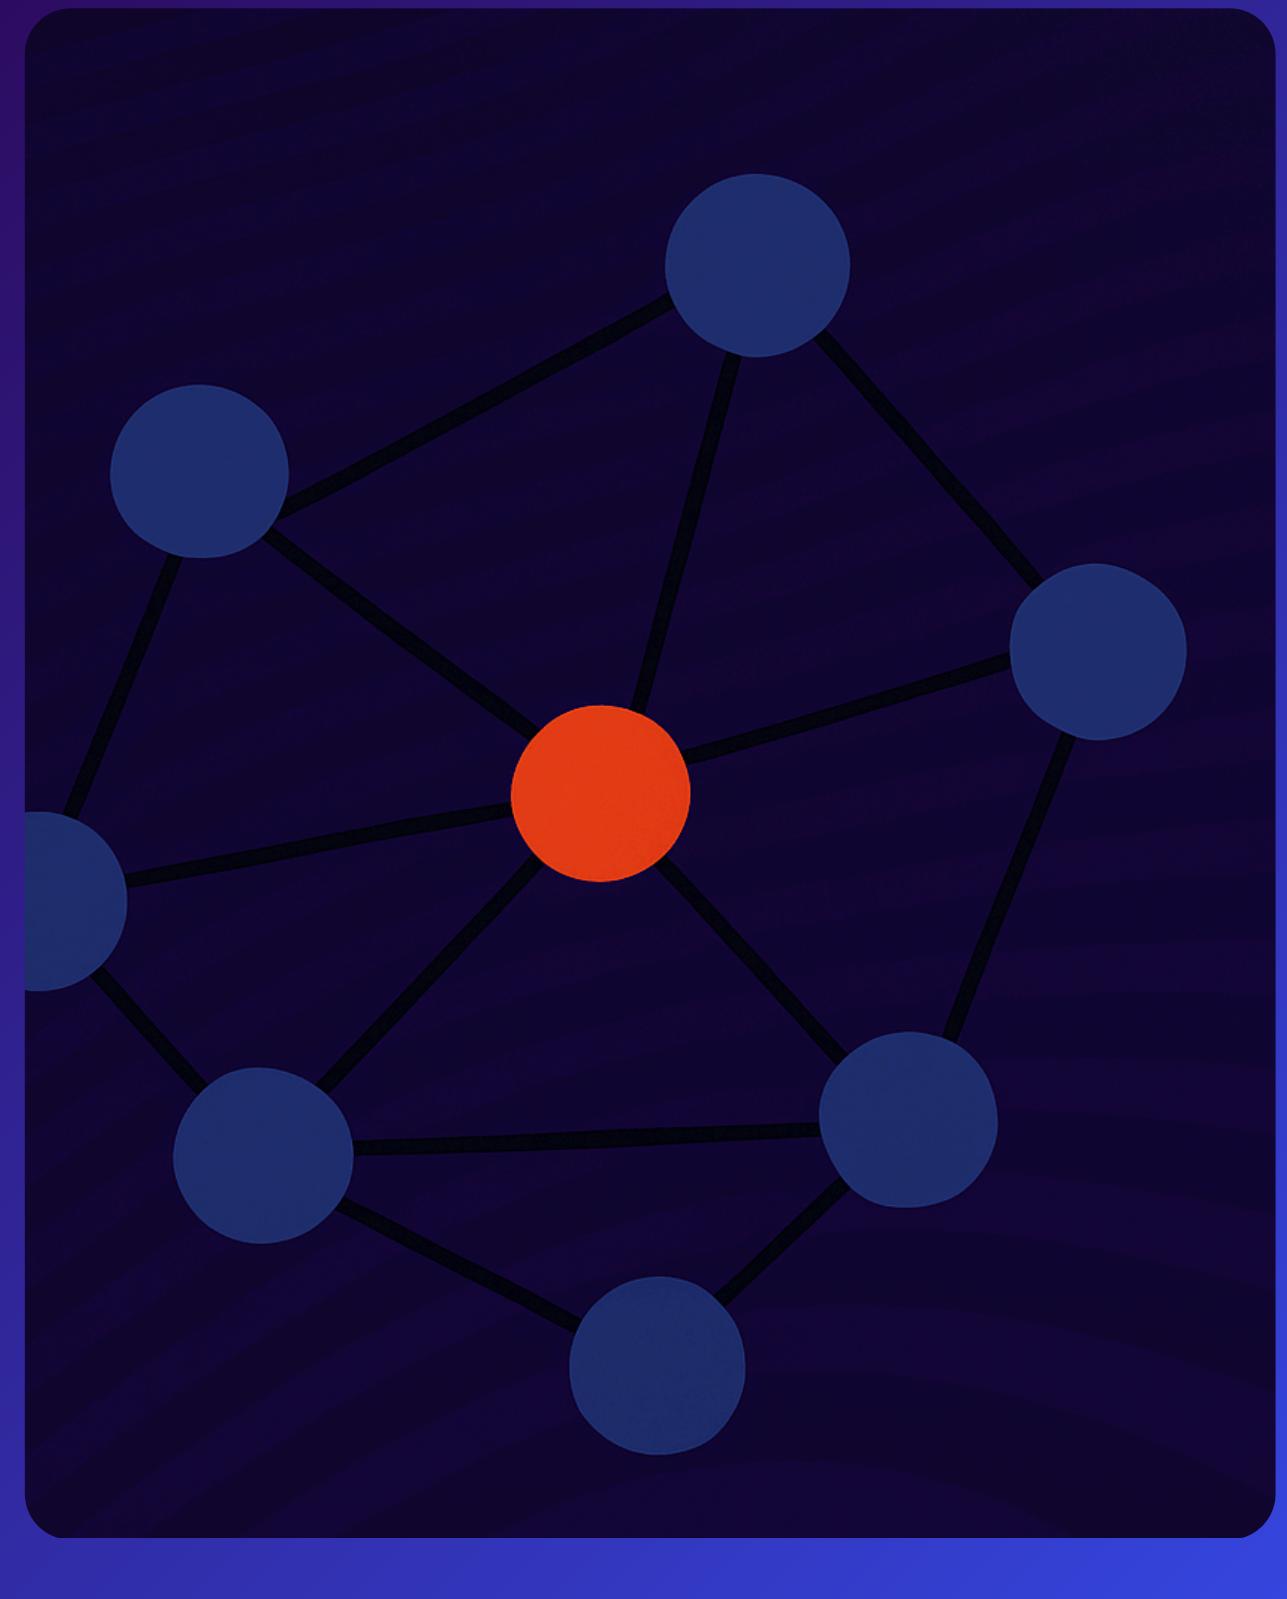
Efficient SSSP Updates via Localized Parallelism

- Generic framework to efficiently update single-source shortest paths in dynamic graphs.
- Avoids full recomputation by updating only affected areas.

Supports both Insertions and Deletions

- Handles edge insertions, deletions, and weight changes.
- Maintains correctness with selective updates.

KEY CONTRIBUTION



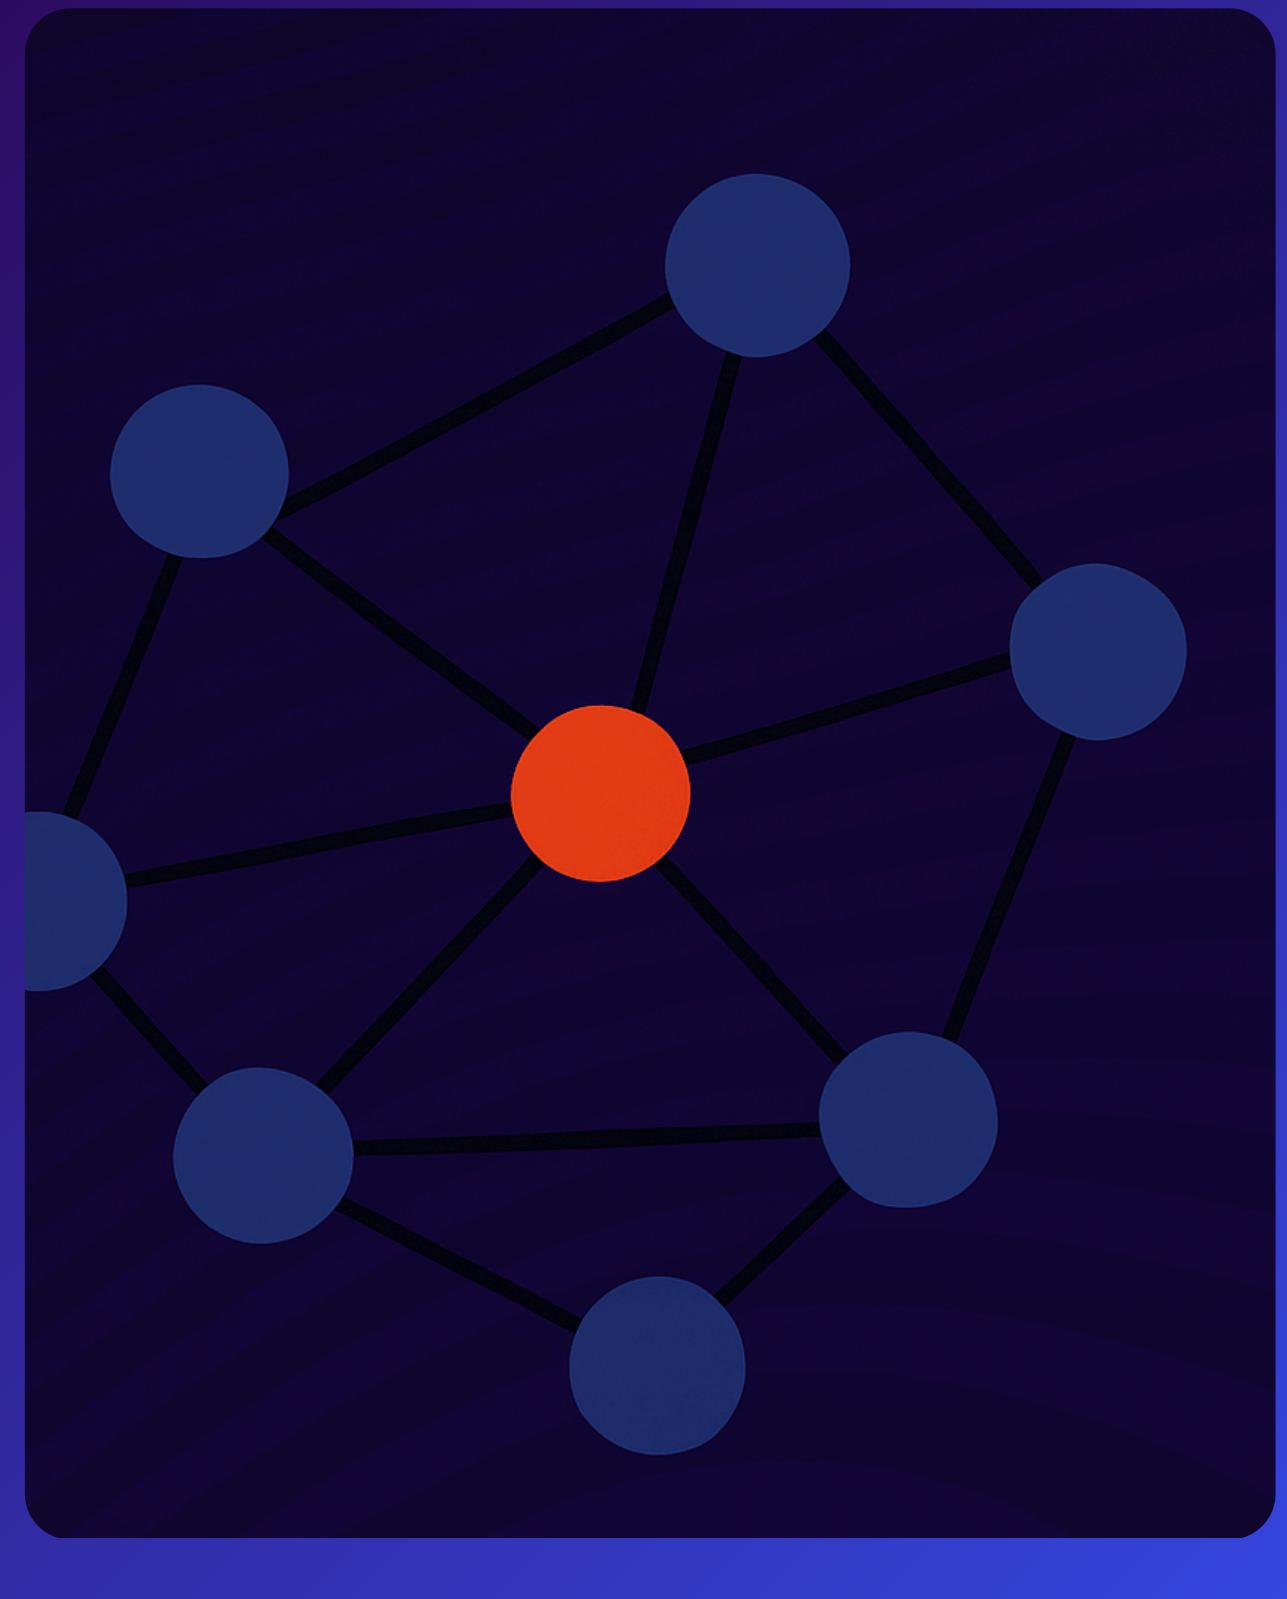
Work-Efficient & Parallel-Friendly

- Minimizes synchronization by using lock-free, iterative updates
- Uses dynamic scheduling to handle load imbalance efficiently

Modular & Flexible Design

- Compatible with multiple SSSP algorithms (e.g., Dijkstra, Bellman-Ford).
- Easily adaptable for different graph types and systems.

KEY CONTRIBUTION



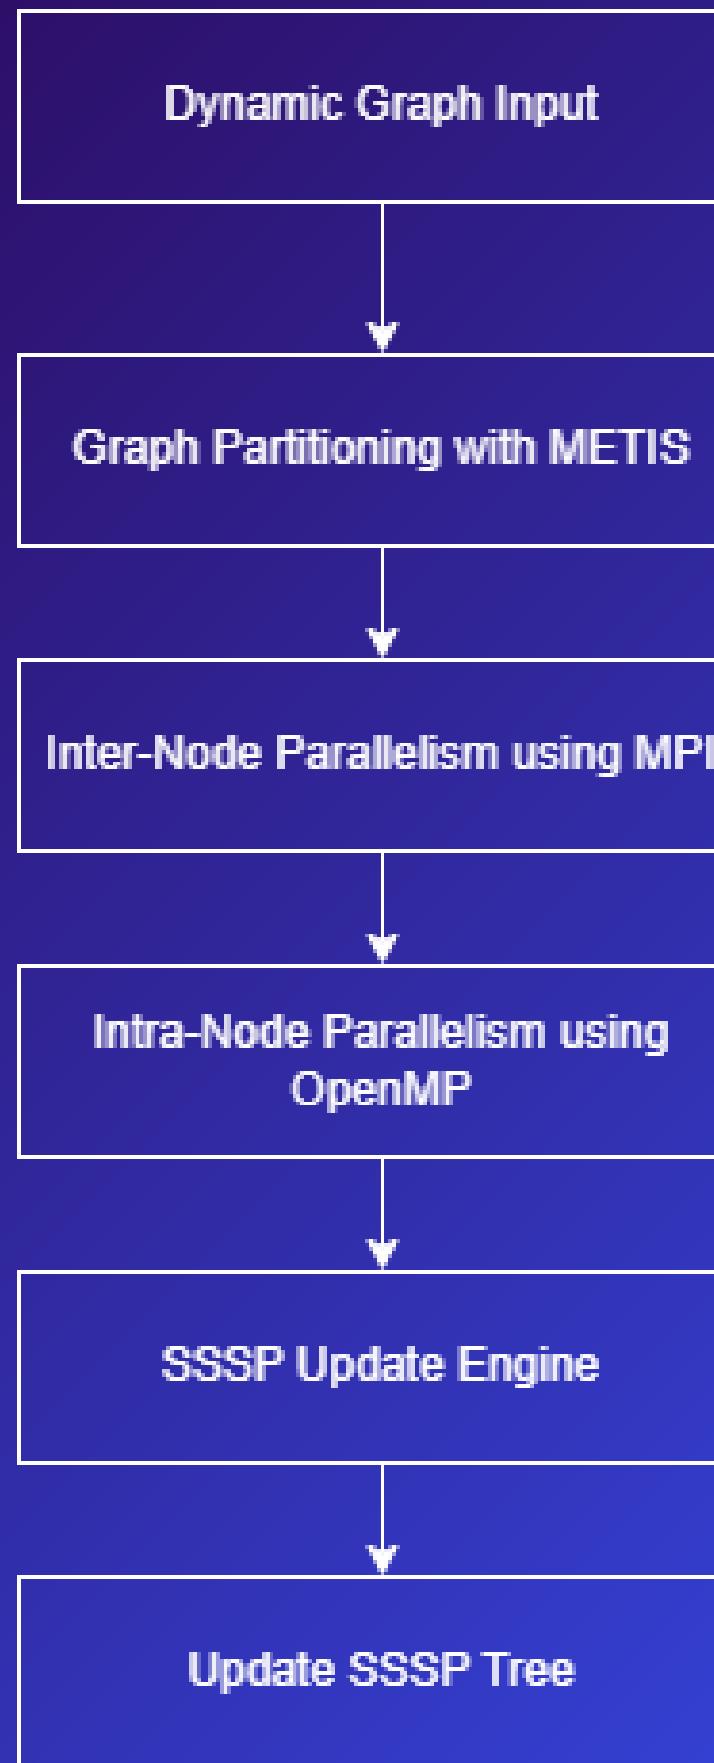
Strong Empirical Results

- Tested on large real-world networks.
- Achieves 10–20× speedup on multi-core setups.

Scalable for Large Dynamic Graphs

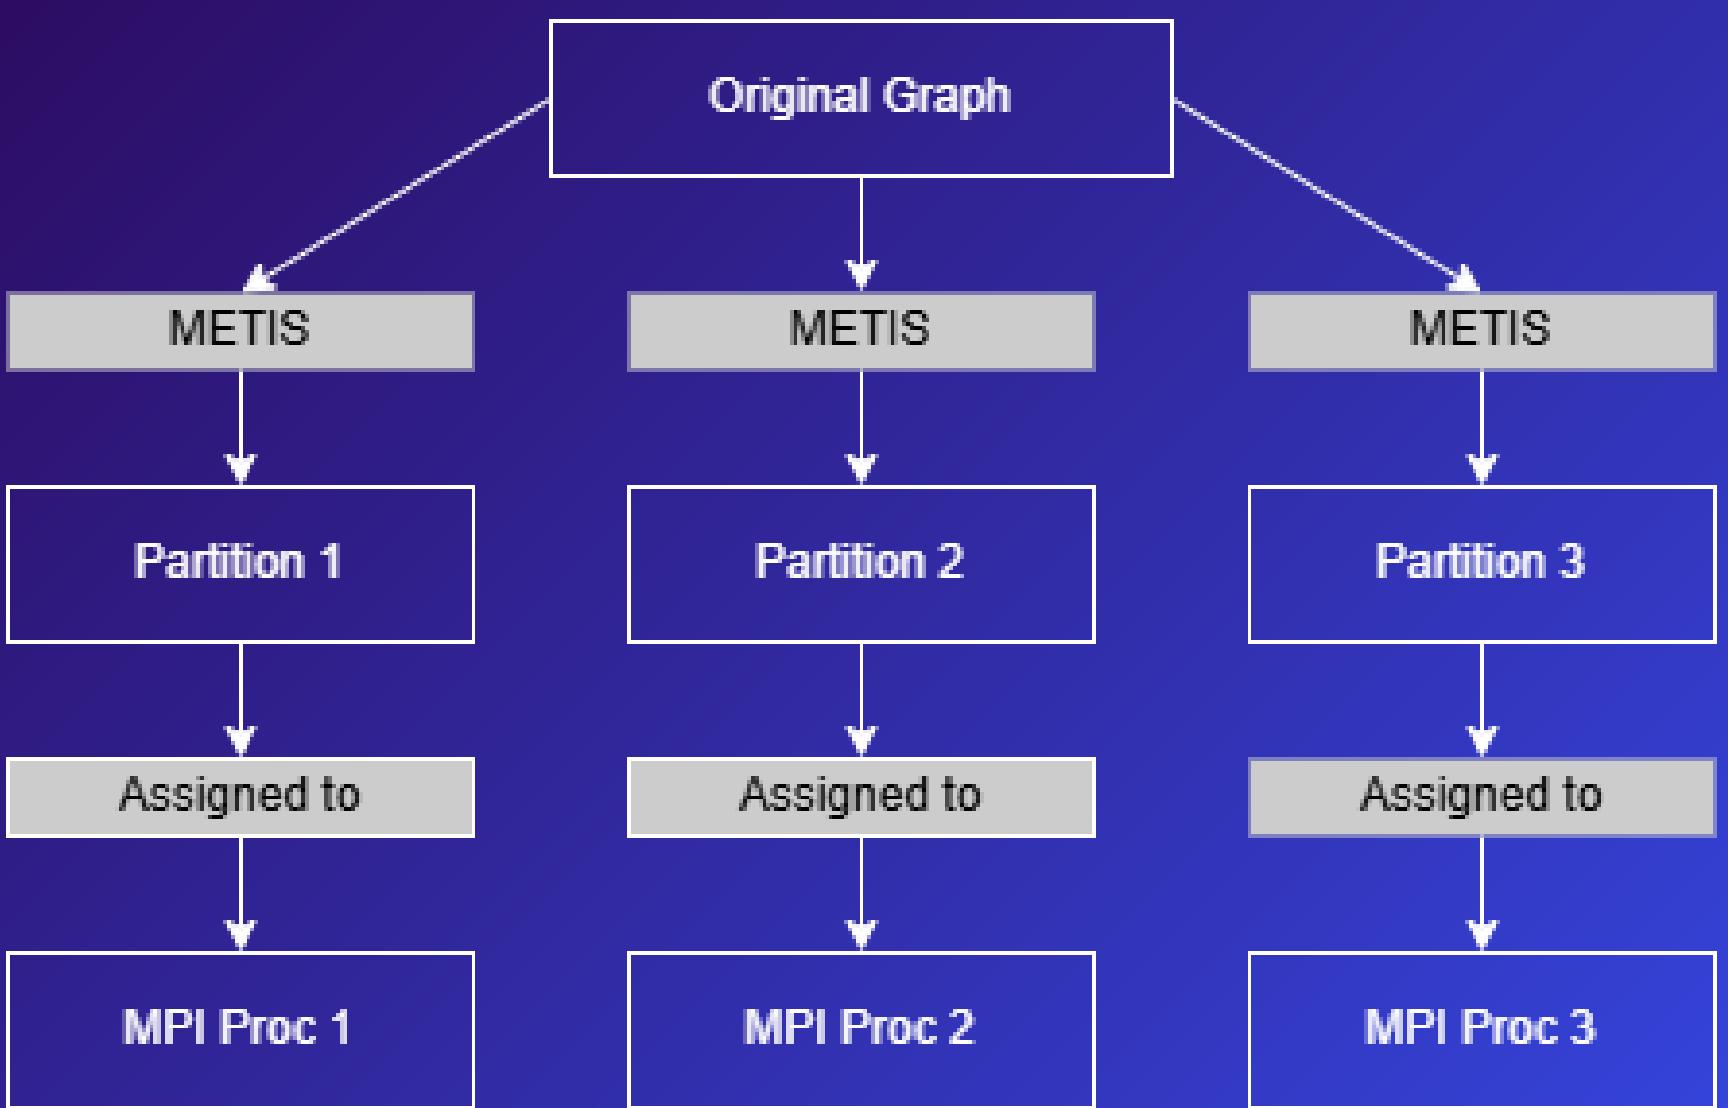
- Efficiently handles large-scale and frequently changing networks.
- Ideal for real-time applications like traffic, social, or web graphs.

PARALLELIZATION STRATEGY



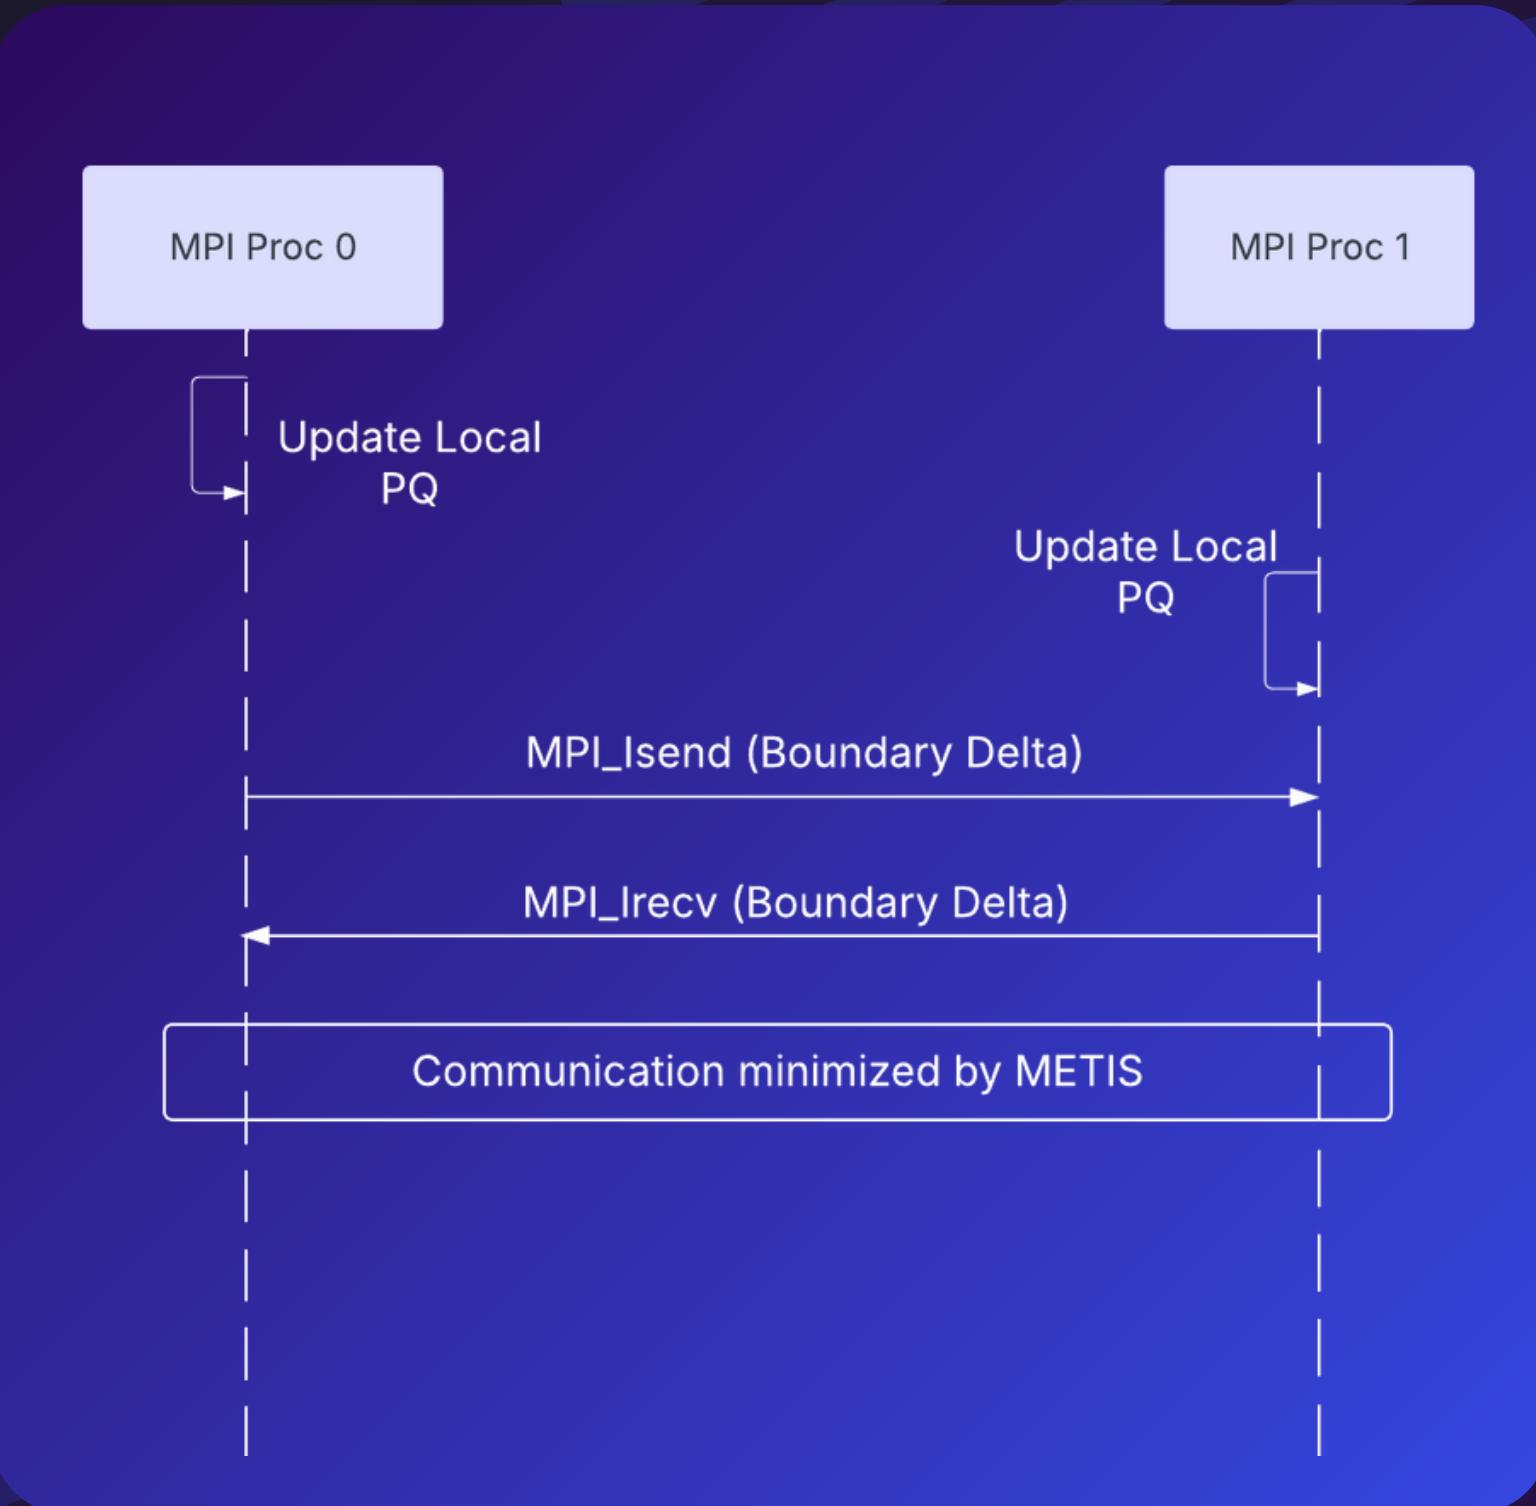
- Graph Partitioning With METIS
- Inter-Node Parallelism with MPI
- Intra-Node Parallelism with OpenMP

GRAPH PARTITIONING WITH METIS



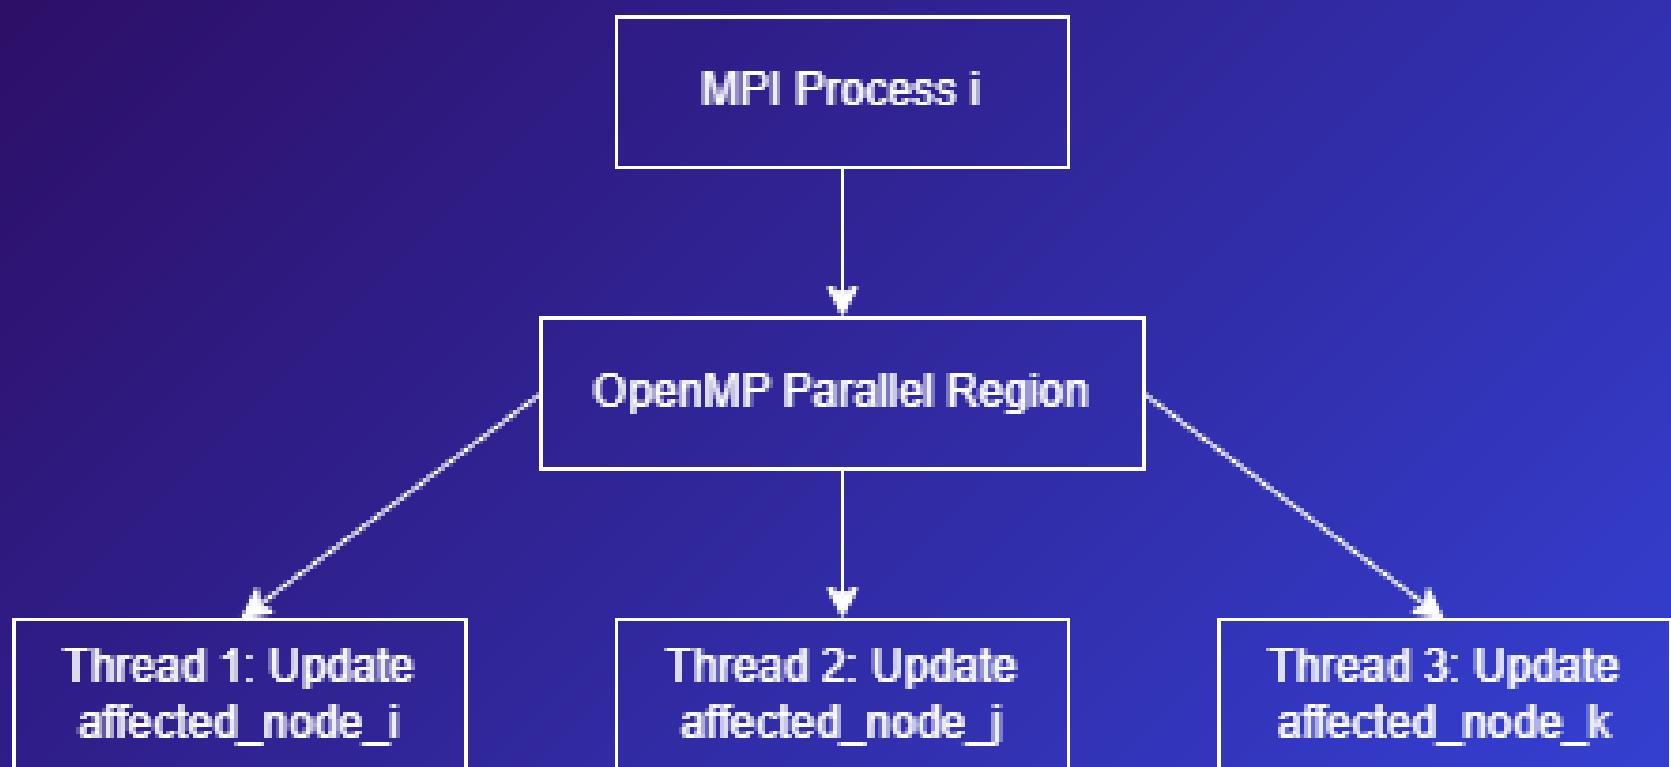
- Input graph is divided into K balanced subgraphs
- Partitions aim to:
 - Balance workload across processes
 - Minimize edge cuts (inter-process communication)
- Each partition is assigned to a separate MPI process

INTER-NODE PARALLELISM-MPI



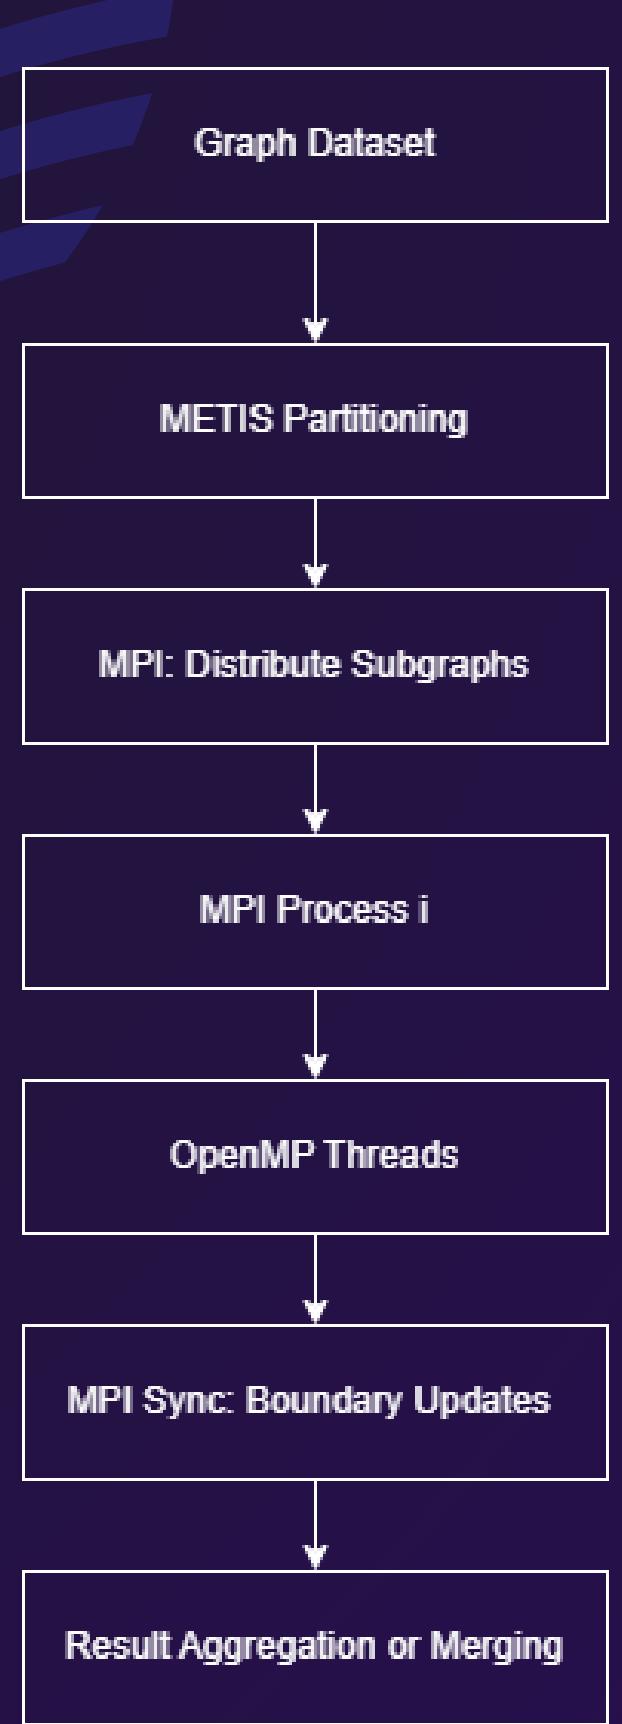
- Each MPI process operates independently on its subgraph
- Communicates updates at partition boundaries
- Uses non-blocking communication to reduce wait times

INTRA-NODE PARALLELISM-OPENMP



- Each MPI process runs multi-threaded logic
- Applies OpenMP for:
 - Processing affected nodes in parallel
 - Relaxing distances
 - Neighbor updates

ALGORITHM FLOW



OPTIMIZATION



1. Reduce Communication:

- send only delta updates

2. Overlap Compute & Comm:

- MPI_Isend is non-blocking

3. Load Balancing:

- initial balance by METIS, monitor runtime skew

4. Thread Scaling:

- use OpenMP dynamic scheduling

PERFORMANCE METRICS WE'LL MEASURE

1. Execution Time:

- Sequential vs MPI vs MPI+OpenMP

2. Speedup:

- Speedup= T_1/T_p

3. Scalability:

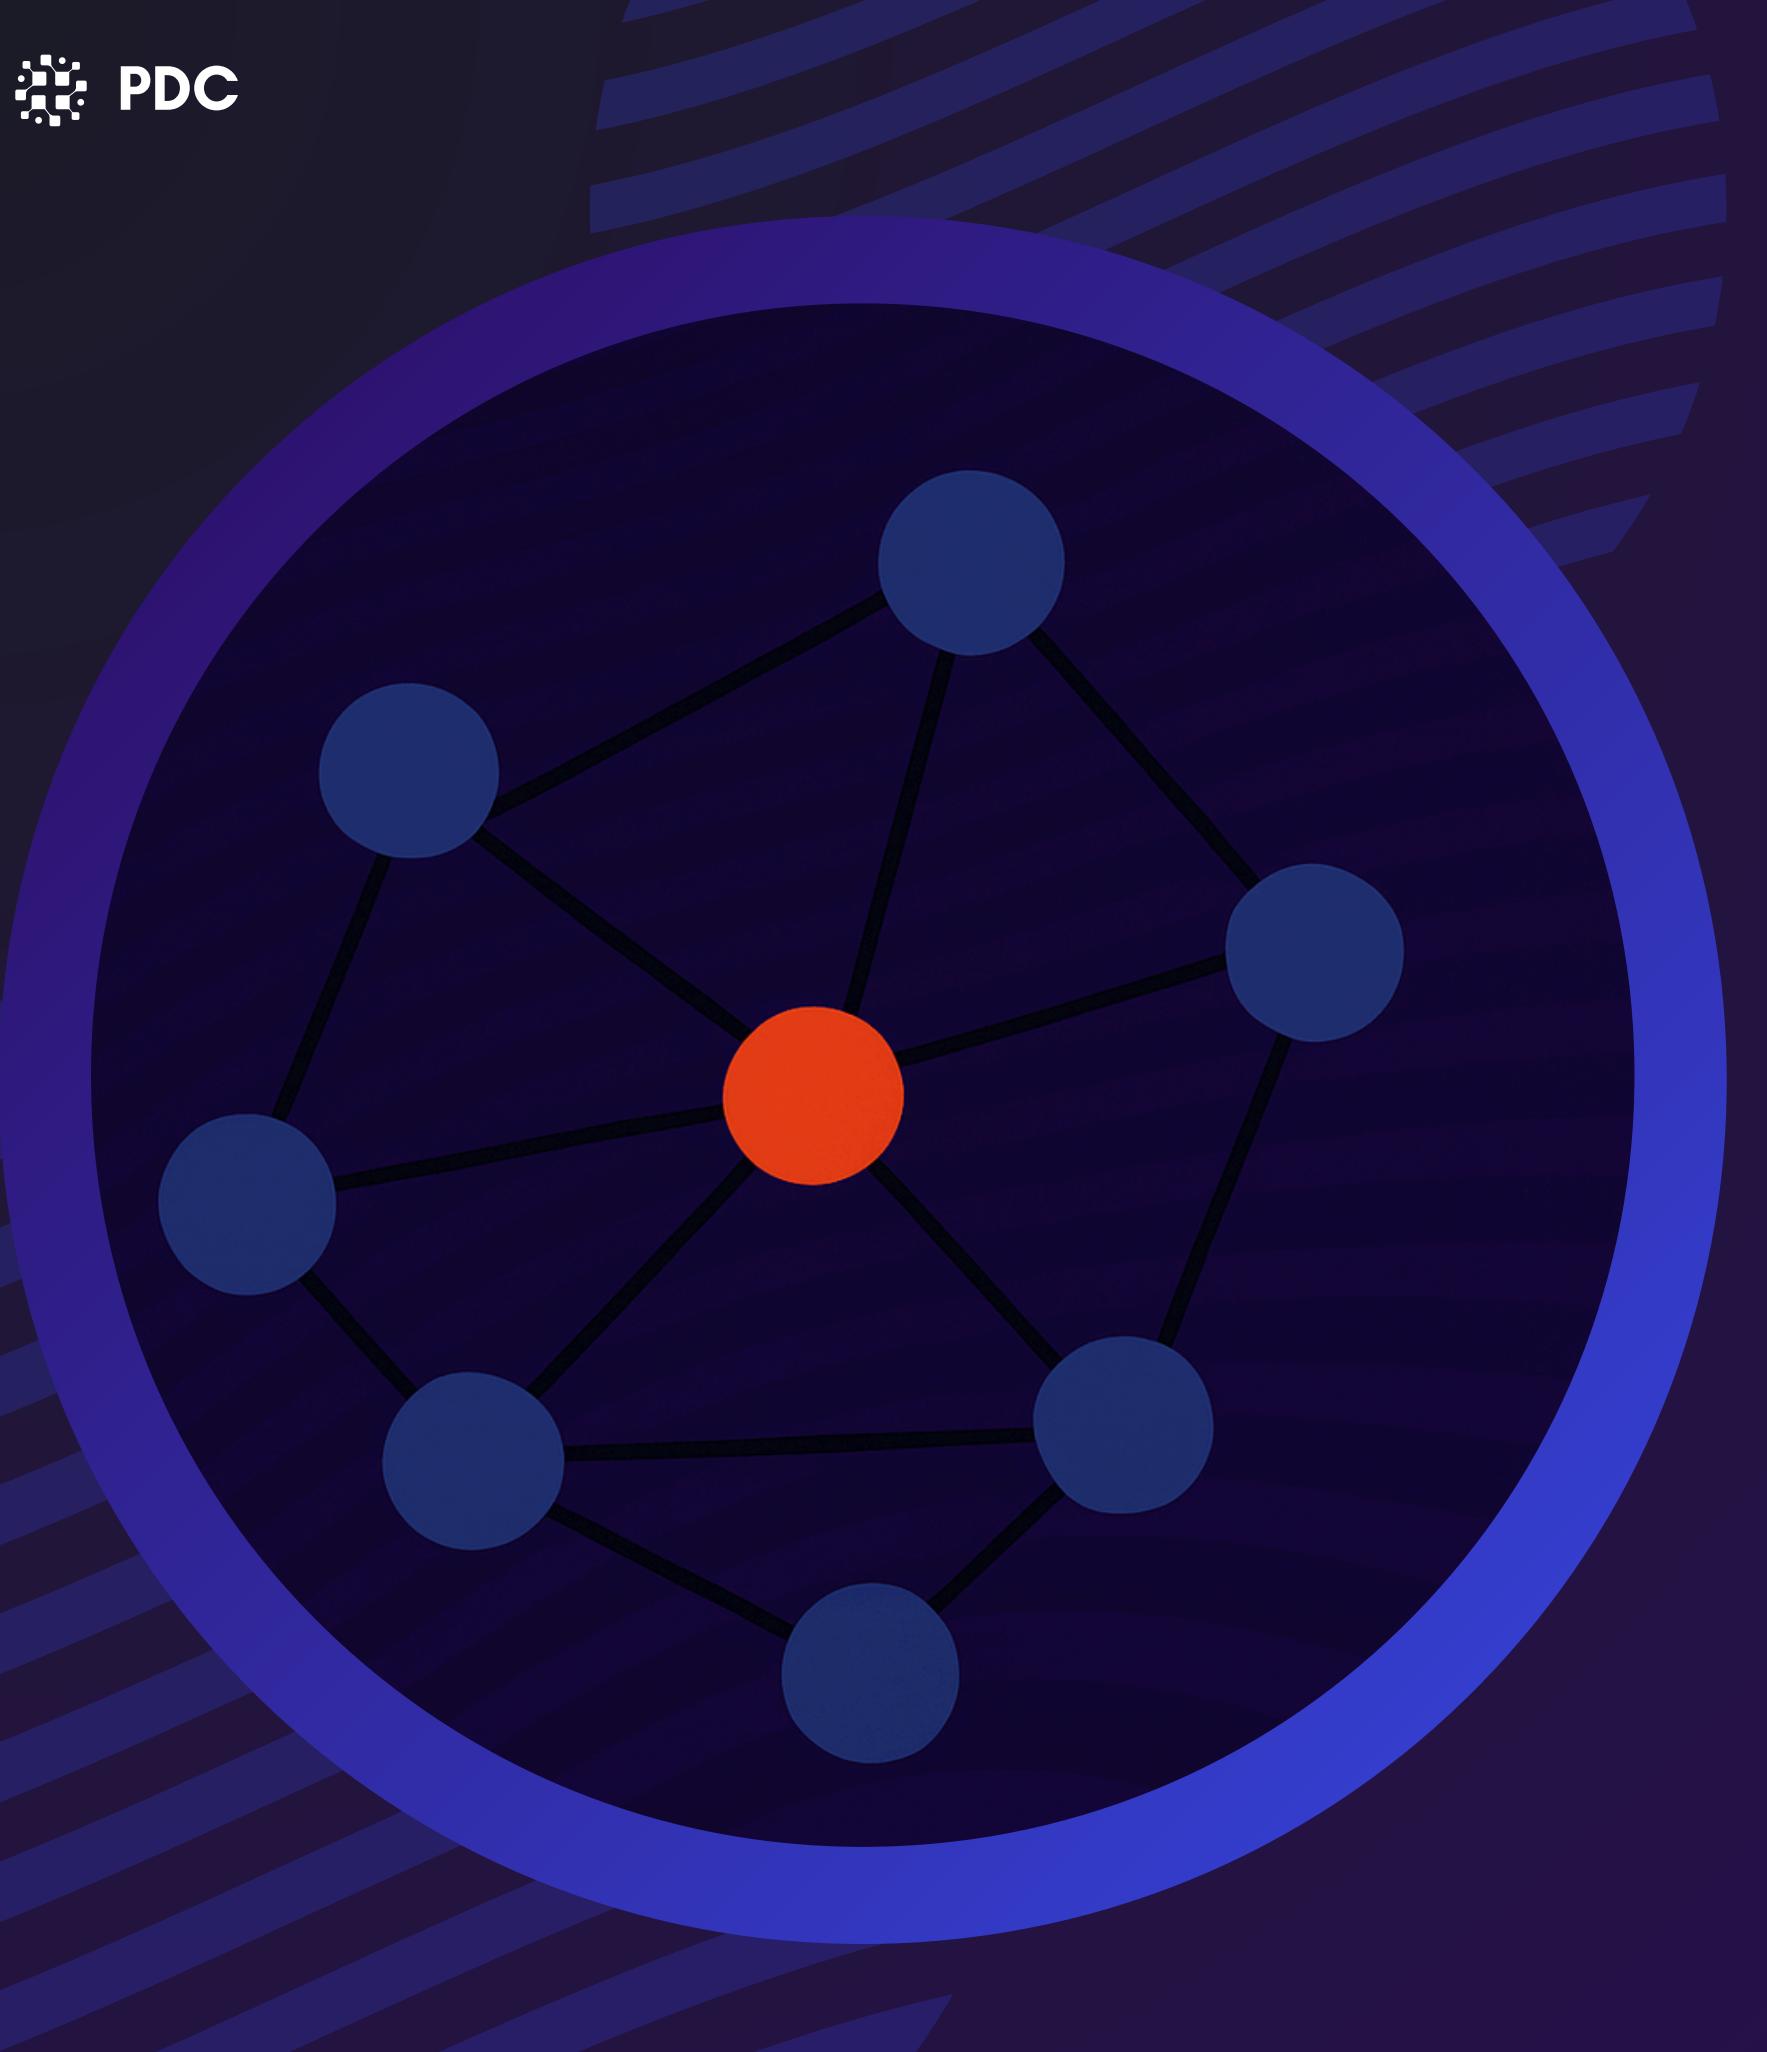
- Strong: fixed graph, more cores
- Weak: graph grows with cores

4. Comm Overhead:

- Time spent in MPI

5. Partitioning Quality:

- Edge cuts and load imbalance



Conclusion